Jiamin Li

Pronouns: she/her/hers ⋄ Email: jiaminli.icy@gmail.com ⋄ Tel: +852 54847038

Ph.D. Candidate (4th year) ⋄ Computer Science ⋄ City University of Hong Kong

RESEARCH INTERESTS

I am primarily interested in Machine Learning systems (MLSys) with a specific focus on the following topics:

- Distributed training and inference: Accelerating large-scale distributed deep neural network (DNN) tasks

- Adaptive and sparse computation: Exploring new computing paradigms to effectively scale DNN models

- Resource scheduling in GPU clusters: Designing efficient scheduling algorithms for DNN tasks

- Simulation: Building accurate performance simulator of DNN training workloads

At present, I am actively engaged in constructing efficient systems tailored for large language models (LLMs).

EDUCATION

The University of Texas at Austin, Austin, TX, United States Apr. 2023 – Present
Visiting Researcher at UTNS Group, Department of Computer Science.
Supervisor: Prof. Aditya Akella
Project: Building Efficient Systems for Large Language Models

City University of Hong Kong, Kowloon, Hong Kong Sep. 2019 – Present
Ph.D. Candidate, Department of Computer Science.
Supervisors: Prof. Hong Xu (The Chinese University of Hong Kong), Prof. Cong Wang (CityU)
Dissertation Title: TBD

City University of Hong Kong, Kowloon, Hong Kong Aug. 2015 – Jul. 2019
B.S., Department of Computer Science (First Class Honours).
Advisor: Dr. Shiqi Wang
Dissertation Title: "Mobile Face Anti-spoofing with Deep Learning"

University of Missouri, Columbia, MO, United States Jul. 2017 – Aug. 2017
Big Data Analysis Summer Program, Department of Computer Science.

WORKING EXPERIENCES

Research Intern, MLSys - AI Lab, ByteDance May. 2019 – May. 2021
Supervisor: Dr. Yibo Zhu
We build a scheduler for training and inference GPU clusters. The key idea is to exploit cluster-level elasticity by loaning idle inferences servers for training and job-level elasticity by scaling jobs to better utilize the dynamic resource pool.

Part-time Research Assistant, City University of Hong Kong May. 2018 – May. 2019
Supervisor: Prof. Hong Xu
During the backpropagation of DNN training, some DNN model gradients exhibit minimal updates across multiple iterations. To address this, we propose an adaptive threshold approach that optimizes communication operations by selectively transmitting only the relevant gradients, effectively reducing data size.

Backend Developer Intern, Jardine Matheson & Co. Limited May. 2017 – May. 2018
Design and develop web services to facilitate employee recruitment for the Group Human Resources department.

PUBLICATIONS

Preprints

[P1] **Jiamin Li**, Cheng Luo, Ziyue Yang, Lei Qu, Peng Cheng, Cong Wang, Hong Xu, "Merak: An Analytical Performance Simulator for Large Scale Distributed Training", under revision.

Conference Proceedings

[C3] **Jiamin Li**, Yimin Jiang, Yibo Zhu, Cong Wang, Hong Xu, " Accelerating Distributed MoE Training and Inference with Lina", USENIX Annual Technical Conference (**ATC**), 2023. (Acceptance Rate = 18.4%)

[C2] **Jiamin Li**, Hong Xu, Yibo Zhu, Zherui Liu, Chuanxiong Guo, Cong Wang, "Lyra: Elastic Cluster Scheduling for Deep Learning", ACM European Conference on Computer Systems (**EuroSys**), 2023. (Acceptance Rate = 16.1%)

[C1] Kaiwei Mo, Chen Chen, **Jiamin Li**, Hong Xu, Chun Jason Xue, "Two-Dimensional Learning Rate Decay: Towards Accurate Federated Learning with Non-IID Data", IEEE International Joint Conference on Neural Networks (**IJCNN**), 2021.

Journals

[J2] Libin Liu, Hong Xu, Zhixiong Niu, Jingzong Li, Wei Zhang, Peng Wang, **Jiamin Li**, Jason Xue Chun, Cong Wang, "ScaleFlux: Efficient Stateful Scaling in NFV", IEEE Transactions on Parallel and Distributed Systems, 2022 (**TPDS**).

[J1] Libin Liu, Chengxi Gao, Peng Wang, Hongming Huang, **Jiamin Li**, Hong Xu, Wei Zhang, "Bottleneck-Aware Non-Clairvoyant Coflow Scheduling with Fai", IEEE Transactions on Cloud Computing, 2021 (**TCC**).

SELECTED RESEARCH PROJECTS

Performance Simulator for Large-scale Distributed DNN Training   Ongoing
Collaborate with Networking Research Group at Microsoft Research Asia
We design a performance simulator to accurately predict the step time of large-scale distributed DNN training tasks. Merak introduces an analytical formulation to compute the all-reduce kernel running time. We also build an ML model to predict the running time slowdown caused by concurrent execution in wait-free backpropagation.

Accelerating Distributed Training and Inference of Mixture-of-Experts models   Completed
Collaborate with MLSys team at ByteDance
MoE models is prone to low efficiency, mainly due to the interleaved all-to-all communication during model computation. In Lina, we propose a communication scheme that prioritizes the all-to-all operations over other communication for MoE training. Besides, we introduce a resource scheduler that exploits expert selection pattern to amid the highly skewed expert popularity.

Elastic Cluster Scheduler for DNN training jobs   Completed
Collaborate with MLSys team at ByteDance
Organizations often build separate training and inference clusters for deep learning. This leads to low utilization in inference clusters and long queuing for jobs in training clusters. Lyra is designed and built to address these problems by cluster-level capacity loaning and job-level elastic scaling.

AWARDS

ATC Student Travel Grant	USENIX ATC	2023
ML and Systems Rising Stars Workshop Program	ML Commons	2023
Research Activity Funds	City University of Hong Kong	2023
EuroSys Student Travel Grant	ACM EuroSys	2023
Full Postgraduate Studentship	City University of Hong Kong	2019 – 2023
Dean's List (College of Engineering)	City University of Hong Kong	2015 – 2019

TEACHING ASSISTANT

2022 Fall, 2019 Fall	CS2311, Computer Programming
2022 Spring, 2020 Spring	CS4296 & CS5296, Cloud Computing
2021 Fall	CS4394 & CS5294, Information Security and Management
2021 Spring	CS4293 & CS6290, Topics on Computer Security
2020 Fall	CS5222, Computer Networks and Internets

PROFESSIONAL SERVICES

Artifact Evaluation Committee	MLSys 2023, ACM EuroSys 2023, ACM CoNEXT 2022, USENIX OSDI 2023, USENIX ATC 2023
Technical Program Committee	IEEE IJCNN 2023
Reviewer	IEEE/ACM Transactions on Networking

TECHNICAL SKILLS

Programming Languages	C++, Python, Bash, Go, LaTeX
Operating Systems	Linux/UNIX
Machine Learning Systems	PyTorch, MXNet, DeepSpeed, HuggingFace, NCCL, CUDA